Long Document Stress Test

1. Section heading number 1

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 1.
- Point two for section 1, somewhat longer so that it wraps onto a second line in most column widths.

2. Section heading number 2

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 2.
- Point two for section 2, somewhat longer so that it wraps onto a second line in most column widths.

3. Section heading number 3

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 3.
- Point two for section 3, somewhat longer so that it wraps onto a second line in most column widths.

4. Section heading number 4

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it.

The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 4.
  - Point two for section 4, somewhat longer so that it wraps onto a second line in most column widths.

5. Section heading number 5

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 5.
  - Point two for section 5, somewhat longer so that it wraps onto a second line in most column widths.

6. Section heading number 6

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 6.
  - Point two for section 6, somewhat longer so that it wraps onto a second line in most column widths.

7. Section heading number 7

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 7.
  - Point two for section 7, somewhat longer so that it wraps onto a second line in most column widths.

8. Section heading number 8

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 8.
- Point two for section 8, somewhat longer so that it wraps onto a second line in most column widths.

9. Section heading number 9

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 9.
- Point two for section 9, somewhat longer so that it wraps onto a second line in most column widths.

10. Section heading number 10

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 10.
- Point two for section 10, somewhat longer so that it wraps onto a second line in most column widths.

11. Section heading number 11

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 11.
- Point two for section 11, somewhat longer so that it wraps onto a second line in most column widths.

12. Section heading number 12

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 12.
- Point two for section 12, somewhat longer so that it wraps onto a second line in most column widths.

13. Section heading number 13

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 13.
- Point two for section 13, somewhat longer so that it wraps onto a second line in most column widths.

14. Section heading number 14

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 14.
- Point two for section 14, somewhat longer so that it wraps onto a second line in most column widths.

15. Section heading number 15

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 15.
- Point two for section 15, somewhat longer so that it wraps onto a second line in most column widths.

16. Section heading number 16

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 16.
- Point two for section 16, somewhat longer so that it wraps onto a second line in most column widths.

17. Section heading number 17

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 17.
- Point two for section 17, somewhat longer so that it wraps onto a second line in most column widths.

18. Section heading number 18

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 18.
- Point two for section 18, somewhat longer so that it wraps onto a second line in most column widths.

19. Section heading number 19

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 19.

- Point two for section 19, somewhat longer so that it wraps onto a second line in most column widths.

20. Section heading number 20

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 20.
- Point two for section 20, somewhat longer so that it wraps onto a second line in most column widths.

21. Section heading number 21

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 21.
- Point two for section 21, somewhat longer so that it wraps onto a second line in most column widths.

22. Section heading number 22

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 22.
- Point two for section 22, somewhat longer so that it wraps onto a second line in most column widths.

23. Section heading number 23

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 23.
- Point two for section 23, somewhat longer so that it wraps onto a second line in most column widths.

24. Section heading number 24

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 24.
- Point two for section 24, somewhat longer so that it wraps onto a second line in most column widths.

25. Section heading number 25

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

- Point one for section 25.
- Point two for section 25, somewhat longer so that it wraps onto a second line in most column widths.