Structural Chunking for Long-Context Retrieval

A. Researcher[1], B. Coauthor[2]
1. Institute of Things 2. Other Lab

Abstract. *Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it. The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.*

1 Introduction

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it.

The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

2 Method

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it.

We define the boundary score *s(i)* as a weighted sum of heading depth and sentence terminality[3].

The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

3 Results

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it.

The mitigation is unglamorous. Chunk boundaries must respect document structure rather than token counts, and the reranker has to be trained on the same distribution it will serve. Everything else is tuning.

4 Discussion

Retrieval quality degrades non-linearly as the corpus grows past the point where the embedding model was calibrated, and the failure is quiet: recall stays flat while precision collapses in the tail. Teams that measure only average relevance will not see it.

3. Terminality is estimated with a small classifier; see supplementary material.